

# Whisper Transcriber

---

A modern, web-based audio transcription tool powered by OpenAI Whisper with real-time progress tracking, GPU acceleration support, and timestamped outputs.

## Features

- **Drag & Drop Upload:** Intuitive web interface for audio file uploads
- **Multiple Whisper Models:** Support for tiny, base, small, medium, large, and v2/v3 variants
- **Real-time Progress:** Chunk-based progress tracking with ETA calculations
- **GPU Acceleration:** Automatic GPU detection with optimized performance estimates
- **Timestamped Transcripts:** Output includes both plain text and timestamped segments
- **Multi-job Dashboard:** Track multiple transcription jobs simultaneously
- **Audio Preprocessing:** Intelligent caching with .npy format for faster processing
- **Flexible Configuration:** Adjustable decoding parameters (temperature, beam size, etc.)
- **Production Ready:** Background processing with cancellation support

## Quick Start

```
# Clone the repository
git clone https://github.com/gopichandbusam/whisper-transcriber.git
cd whisper-transcriber

# Make management script executable
chmod +x manage.sh

# Setup environment and dependencies
./manage.sh setup

# Run the application
./manage.sh run
```

Visit <http://127.0.0.1:5000> in your browser and start transcribing!

## Requirements

- Python 3.9+
- macOS or Linux
- Sufficient disk space for model weights (100MB - 6GB depending on model)
- Optional: CUDA-compatible GPU for acceleration

## Configuration

### Model Selection

```
# Use a specific model size
WHISPER_MODEL_SIZE=small ./manage.sh run

# Enable chunk-based progress for better tracking
PROGRESS_MODE=chunks ./manage.sh run

# Custom port
PORT=8080 ./manage.sh run
```

## Available Models

Model	Speed	Accuracy	RAM Usage	Best For
tiny	⚡ ⚡ ⚡ ⚡	★	~1 GB	Quick previews
base	⚡ ⚡ ⚡	★★	~1.3 GB	General use
small	⚡ ⚡	★★★	~2.5 GB	Balanced quality
medium	⚡	★★★★	~5 GB	High quality
large	🐌	★★★★★	~6+ GB	Maximum accuracy

## 📖 Documentation

- **Setup Guide:** See main README sections for detailed installation
- **API Reference:** Check [/health](#) endpoint for runtime configuration
- **Troubleshooting:** Common issues and solutions in main README

## 🤝 Contributing

We welcome contributions! Here's how you can help:

1. **Fork** the repository
2. **Create** a feature branch (`git checkout -b feature/amazing-feature`)
3. **Commit** your changes (`git commit -m 'Add amazing feature'`)
4. **Push** to the branch (`git push origin feature/amazing-feature`)
5. **Open** a Pull Request

## Development Setup

```
./manage.sh shell # Activate development environment
python app.py      # Run in debug mode
```

## 🐛 Issues & Support

- **Bug Reports:** [Open an issue](#)

- **Feature Requests:** [Start a discussion](#)
- **Questions:** Check existing issues or create a new one

## License

This project is licensed under the MIT License - see the [LICENSE](#) file for details.

## Acknowledgments

- [OpenAI Whisper](#) for the amazing speech recognition model
- [Flask](#) for the web framework
- All contributors who help improve this project

## Project Stats

 Stars repo not found

 Forks repo not found

issues repo not found

license repo not found

---

Made with ❤️ by [Gopichand Busam](#)

If you find this project helpful, please consider giving it a ⭐ !